

Log Analysis in Linux and Splunk

Introduction

Objective

Log analysis is one of the most widely used skills for an analyst. Logs can give indicators of what might be happening to systems on the network. Logs can indicate precursors to compromise, can contain information about a compromise, and also can contain information relevant to post exploitation activities. There are various tools that can be used to analyze logs. There are commercial tools like Splunk as well as free tools like awk, gawk, and grep. The important thing is to be able to look at the logs and parse the relevant information that you're looking for regardless of the tool used.

Overview

Splunk is a widely used commercial log aggregation tool. It is a great tool for ingesting data and then being able to help you analyze network incidents. There are other ways to view log files besides using Splunk. For instance, using grep, gawk, and awk can provide you with similar log parsing results, but those are more arduous methods and those tools are more command line based and require the analyst to remember many various commands and switches.

Outcomes

In this lab, you will learn to:

- 1 Use Linux commands to search Linux logs.
- 2 Use Splunk to analyze network traffic and logs.

Key Term	Description
cat	A Linux command used to show the output of data

Key Term	Description
Splunk	A commercial tool with the ability to analyze large log files
grep (global regular expressions print)	A Linux tool to parse information
Leafpad	A GUI text editor program for Linux (similar to Windows Notepad)
tail	The command to show the last few lines of a file in Linux

Reading Assignment


Introduction

Log analysis is one of the most widely used skills for an analyst. Logs can give indicators of what might be happening to systems on the network. Logs can:

- Indicate precursors to compromise
- Contain information about a compromise
- Contain information relevant to post exploitation activities

There are various tools that can be used to analyze logs. There are commercial tools like Splunk as well as free tools like the Linux commands `awk`, `gawk`, and `grep`. The important thing is to be able to look at the logs and parse the relevant information that you're looking for regardless of the tool used.

Figure 1 shows the lab topology.



In this lab, you will learn to:

- Use Linux commands to search Linux logs.
- Use Splunk to analyze network traffic and logs.

Linux Log Analysis

Linux keeps a large number of log files that keep track of events on the system. One of the most important log files is the log “secure” which logs security incidents. The secure file is located in /var/log on a Red Hat variant of Linux. The Ubuntu/Debian equivalent to secure is auth.log. This log file tracks the following information:

- Creation of new user accounts
- Creation of new group accounts
- Logging on from a remote system
- Changing a user’s password

Another important file on a Linux system is the access_log file, which keeps track of web server connections. In Linux, the most common web server used is Apache.

The Apache Server keeps records of the connections made to the website, including:

- IP addresses
- User agents
- Date/timestamps

In this lab, you will use various utilities to search through logs in Linux.

cat

The cat command displays the content of files to the screen. It works well with small files but will scroll off the screen for larger files. The command cat is short for concatenate. Figure 2 shows a cat example.

```
root@kali:~# cat ex191112.log
#Software: Microsoft Internet Information Services 7.0
#Version: 1.0
#Date: 2019-11-12 22:22:22
#Fields: time c-ip cs-method cs-uri-stem sc-status sc-win32-status
22:22:22 175.45.176.200 [1]USER administrator 331 0
22:22:22 175.45.176.200 [2]USER administrator 331 0
22:22:22 175.45.176.200 [3]USER administrator 331 0
22:22:22 175.45.176.200 [4]USER administrator 331 0
22:22:22 175.45.176.200 [5]USER administrator 331 0
22:22:22 175.45.176.200 [2]PASS - 530 1326
22:22:22 175.45.176.200 [1]PASS - 530 1326
22:22:22 175.45.176.200 [3]PASS - 530 1326
22:22:22 175.45.176.200 [5]PASS - 530 1326
```

Figure 2 - cat Command

head

The head command displays the top 10 lines of a file by default. Different switches can be used with head to show a different number of lines at the beginning of a file. Figure 3 shows a head example.

```
root@kali:~# head ex191112.log
#Software: Microsoft Internet Information Services 7.0
#Version: 1.0
#Date: 2019-11-12 22:22:22
#Fields: time c-ip cs-method cs-uri-stem sc-status sc-win32-status
22:22:22 175.45.176.200 [1]USER administrator 331 0
22:22:22 175.45.176.200 [2]USER administrator 331 0
22:22:22 175.45.176.200 [3]USER administrator 331 0
22:22:22 175.45.176.200 [4]USER administrator 331 0
22:22:22 175.45.176.200 [5]USER administrator 331 0
22:22:22 175.45.176.200 [2]PASS - 530 1326
```

Figure 3 - head Command

tail

The tail command displays the last 10 lines of a file by default. You can also select the number of lines at the end of a file using the -n option. Figure 4 shows an example of the tail command.

```
root@kali:~# tail ex191112.log
22:25:42 175.45.176.200 [3]USER administrator 331 0
22:25:42 175.45.176.200 [5]PASS - 230 0
22:25:42 175.45.176.200 [2]PASS - 530 1326
22:25:42 175.45.176.200 [5]QUIT - 230 0
22:25:42 175.45.176.200 [2]USER administrator 331 0
22:25:42 175.45.176.200 [4]PASS - 530 1326
22:25:42 175.45.176.200 [3]PASS - 530 1326
22:25:42 175.45.176.200 [1]PASS - 530 1326
22:25:42 175.45.176.200 [2]PASS - 530 1326
```

Figure 4 - tail Command

wc

The wc command displays the number of lines, words, and characters of a file. Figure 5 shows an example of the wc command.

```
root@kali:~# wc -l ex191112.log
361717 ex191112.log
```

Figure 5 - wc Command

Regular Expressions

Regular expressions (RegEx) are defined as a search pattern that allows you to do powerful searches and replacements in strings and files. You also can use regular expressions for input validation in programs and web applications. Regular expressions are one of the most powerful tools in a system administrator's toolbox. Once mastered, system administrators can write very powerful scripts to automate repetitive tasks.

History of Regular Expressions

Regular expressions concept arose in the 1950s by American Mathematician Stephen Cole Keene. Keene developed a regular language. Regular expressions are used in programming languages for input validation, search engines, word processors, text editors, and other applications.

grep

The grep (stands for global regular expressions print) command in Linux is very powerful for indexing text files. The grep command is used to search for a search pattern (regular expression) in a file.

Literal Matching

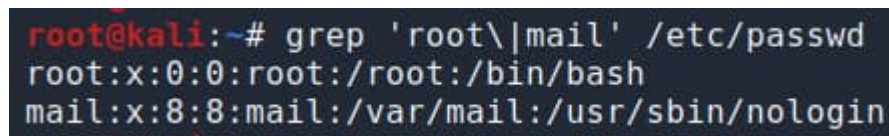
The most basic search using grep is using the following command:

```
grep <search pattern> <file-name>
```

For example, to find the root in the /etc/passwd file, you would use the following command: `grep root /etc/passwd`. You will get back all the instances of root in the lines in the /etc/passwd file. This is what is called literal matching.

Boolean “or”

Let's say you want to search for root or mail. You would use the following command: `grep 'root\|mail' /etc/passwd`. This will return all the instances of root or mail in /etc/passwd. You use `\|` for the or and you put the regular expression in single quotes to tell the command that this is a regular expression. Figure 6 shows an example grep using the Boolean “or.”



```
root@kali:~# grep 'root\|mail' /etc/passwd
root:x:0:0:root:/root:/bin/bash
mail:x:8:8:mail:/var/mail:/usr/sbin/nologin
```

Figure 6 - grep Command

Anchoring

There are meta-characters that tell you where in a line to match. For example, the command `grep '^root' /etc/passwd` will only return items that have root in the beginning of the string. The dollar sign (\$) is another meta-character. With the dollar sign meta-character, grep will match anything at the end of a string. So, the command `grep 'bash$' /etc/passwd` will match strings with bash at the end. Now, if you use `grep '^$' /etc/passwd`, this will match empty lines in a file such as `grep '^$' /etc/passwd`. In the case that there are no empty lines in /etc/passwd, so nothing will return.

Single Character Matching

The period is a meta-character that allows you to match any single character. For example, the command `grep 'r.0t' /etc/passwd` will match any character in the second position of the search pattern where the period is located.

Bracket Expressions

Bracket expressions allows you to match a group of characters by putting them into `[]`. For example, if you want to match all strings that have either a small d or a capital D, you could use `grep '[dD]' /etc/passwd`. This will return all lines with a small d or capital D in them. If you want to search for all lines that begin with a lowercase letter, you would use this command `grep '^[a-z]' /etc/passwd`.

There are special bracket expressions that can be used in your regular expressions.

Quantifier	Character Classes
<code>[:alnum:]</code>	Alphanumeric characters
<code>[:alpha:]</code>	Alphabetic characters
<code>[:blank:]</code>	Space and tab
<code>[:digit:]</code>	Digits
<code>[:lower:]</code>	Lowercase letters
<code>[:upper:]</code>	Uppercase letters

Quantification

The following table shows the different quantifiers, descriptions, and regular expression examples.

Quantifier	Description	Example
<code>*</code>	Matches the preceding item zero or more times	<code>'n*login'</code>
<code>?</code>	Matches the preceding item zero or one time	<code>'n?login'</code>
<code>+</code>	Matches the preceding item one or more times	<code>'no+'</code>
<code>{n}</code>	Matches the preceding item exactly n times	<code>'[:digit:]{3}'</code>
<code>{n,}</code>	Matches the preceding item at least n times	<code>'[:digit:]{9,}'</code>
<code>{,m}</code>	Matches the preceding item at most m times	<code>'[:digit:]{,2}'</code>
<code>{n,m}</code>	Matches the preceding item from n to m times	<code>'[:digit:]{9,12}'</code>

Grouping

Grouping allows you to group patterns together and reference them as a single entity using (). For example, `grep '(no\)' /etc/passwd` will show everything that has no in the word.

Backslash Expressions

The following table shows special backslash expressions and examples.

Expression	Description
<code>\b</code>	Matches a word boundary
<code>\w</code>	Matches a word
<code>\s</code>	Matches a space

For example, `grep '\b[nN]fs\b' /etc/passwd` will return anything that is N or n from the term nfs (network file system).

gawk

The gawk command is a program that does pattern matching, record processing, and other forms of text manipulation.

sort

The sort command will sort a list of values in a file.

uniq

The uniq command removes duplicate adjacent lines from sorted file or from standard input.

Figure 7 shows an example of using gawk, sort, and uniq. The command will display the total number of unique values in the first column of the log file.


```
root@kali:~# cat ex191112.log | gawk '{ print $1}' | sort | uniq -c
1835 22:22:22
1812 22:22:23
1902 22:22:24
1966 22:22:25
1957 22:22:26
1877 22:22:27
1960 22:22:28
1952 22:22:29
1993 22:22:30
1981 22:22:31
1897 22:22:32
1926 22:22:33
1942 22:22:34
1977 22:22:35
2000 22:22:36
1991 22:22:37
1822 22:22:38
1963 22:22:39
1969 22:22:40
1974 22:22:41
1938 22:22:42
1944 22:22:43
```

Figure 7 - cat Command

Splunk

Splunk is a log analysis tool that will allow you to ingest several different types of log formats, including web logs, system logs, event viewer files, and Packet Capture Application (PCAP) data. Splunk is a commercial product, but it has a free edition. With the free edition of Splunk, there is a parse limit of 500 MB per day within a 24-hour time frame. Figure 8 shows the log in screen of Splunk.

Splunk helps “capture, index, and correlate real-time data in a searchable repository from which it can generate graphs, reports, alerts, dashboards, and visualizations.” Source: [Link](#)

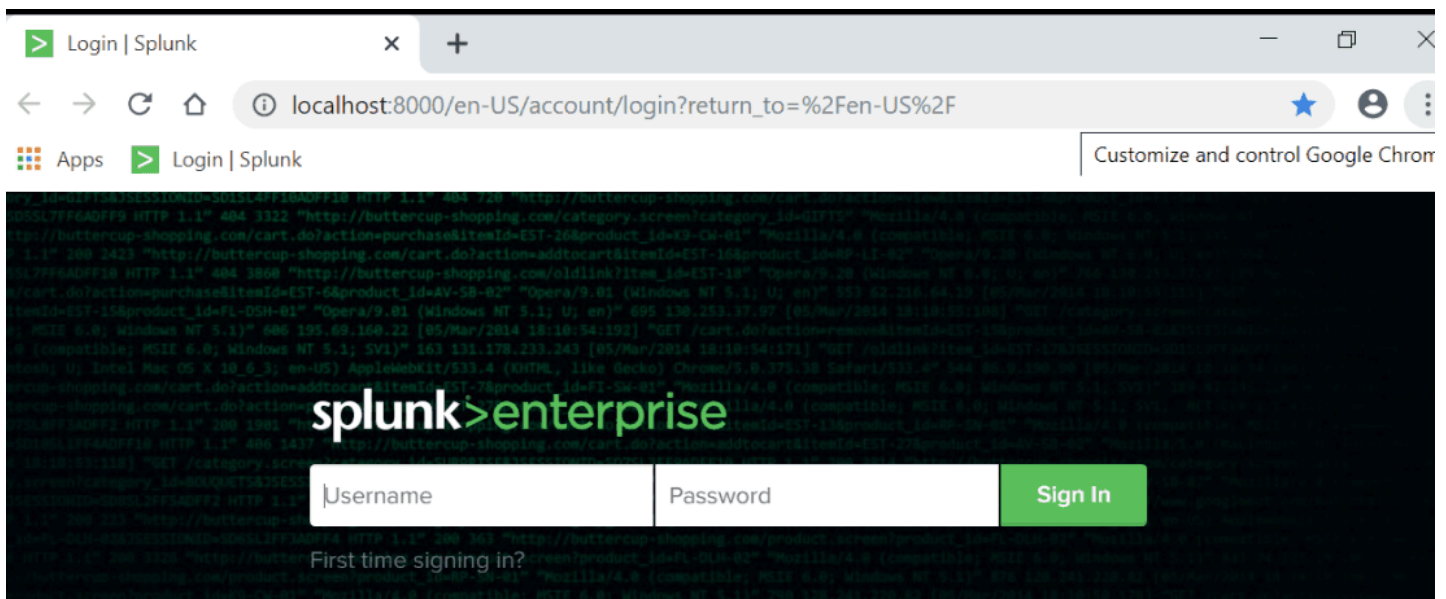


Figure 8 - Splunk

CONCLUSION:

Splunk and Linux command-line log analyses provide the same result, but the Linux results are provided faster. The disadvantage to the Linux method is remembering the use and syntax of the commands. Frequent use and practice with the Linux commands and the available switches will make you more familiar with them. Using Splunk is always another option.

Using the Command Line in Linux to Parse Log files

- 1 **Click** on the **Kali virtual machine**. **Log in** with the **username** of `root` and the **password** of `toor`.



Note: The password of **toor** will not be displayed when you type it for security purposes.

- 2 **Open** a **terminal** on **Kali**. **Type** the following **command** and **press Enter** to get the version of **Linux**.

```
root@kali:~# uname -a
```

- 3 **Type** the following **command** and **press Enter** to list files and folders in the **root** **directory**.

```
root@kali:~# ls
```

- 4 **Type** the following **command** and **press Enter** to view the **log files**.

```
root@kali:~# cat ex191112.log
```

```
root@kali:~# cat ex191112.log
#Software: Microsoft Internet Information Services 7.0
#Version: 1.0
#Date: 2019-11-12 22:22:22
#Fields: time c-ip cs-method cs-uri-stem sc-status sc-win32-status
22:22:22 175.45.176.200 [1]USER administrator 331 0
22:22:22 175.45.176.200 [2]USER administrator 331 0
22:22:22 175.45.176.200 [3]USER administrator 331 0
22:22:22 175.45.176.200 [4]USER administrator 331 0
22:22:22 175.45.176.200 [5]USER administrator 331 0
22:22:22 175.45.176.200 [2]PASS - 530 1326
22:22:22 175.45.176.200 [1]PASS - 530 1326
22:22:22 175.45.176.200 [3]PASS - 530 1326
22:22:22 175.45.176.200 [5]PASS - 530 1326
```



Note: The file is very long and has a lot of lines. We will use the `wc` command in Linux to determine the number of lines.

- 5 **Type** the following **command** and **press Enter** to view the options for `wc`. `wc` is a program that displays the line, word, character, and byte count of a file.

```
root@kali:~# wc --help
```

- 6 **Type** the following **command** and **press Enter** to view the number of lines in the file.

```
root@kali:~# wc -l ex191112.log
```

```
root@kali:~# wc -l ex191112.log
361717 ex191112.log
```



Note: The number to the left of the file indicates that this log file has 361,717 lines. That is a large amount of data.

The `head` command can be used in Linux to display the beginning of a file.

- 7 **Type** the following **command** and **press Enter** to view the options for `head`. `head` is a program that displays first lines of a file.

```
root@kali:~# head --help
```

- 8 **Type** the following **command** and **press Enter** to view the **first 10 lines** of the **log files**.

```
root@kali:~# head ex191112.log
```

```
root@kali:~# head ex191112.log
#Software: Microsoft Internet Information Services 7.0
#Version: 1.0
#Date: 2019-11-12 22:22:22
#Fields: time c-ip cs-method cs-uri-stem sc-status sc-win32-status
22:22:22 175.45.176.200 [1]USER administrator 331 0
22:22:22 175.45.176.200 [2]USER administrator 331 0
22:22:22 175.45.176.200 [3]USER administrator 331 0
22:22:22 175.45.176.200 [4]USER administrator 331 0
22:22:22 175.45.176.200 [5]USER administrator 331 0
22:22:22 175.45.176.200 [2]PASS - 530 1326
```

The tail command can be used in Linux to display the end of a file.

- 9 **Type** the following **command** and **press Enter** to view the options for **tail**. tail is a program that displays a number of lines at the end of a file.

```
root@kali:~# tail --help
```

- 10 **Type** the following **command** and **press Enter** to view the **last 10 lines** of the **log files**.

```
root@kali:~# tail -n 10 ex191112.log
```

```
root@kali:~# tail ex191112.log
22:25:42 175.45.176.200 [3]USER administrator 331 0
22:25:42 175.45.176.200 [5]PASS - 230 0
22:25:42 175.45.176.200 [2]PASS - 530 1326
22:25:42 175.45.176.200 [5]QUIT - 230 0
22:25:42 175.45.176.200 [2]USER administrator 331 0
22:25:42 175.45.176.200 [4]PASS - 530 1326
22:25:42 175.45.176.200 [3]PASS - 530 1326
22:25:42 175.45.176.200 [1]PASS - 530 1326
22:25:42 175.45.176.200 [2]PASS - 530 1326
```

- 11 **Type** the following **command** and **press Enter** to view the **first 25 lines** of the **log files**.

```
root@kali:~# head -n 25 ex191112.log
```

```
root@kali:~# head -n 25 ex191112.log
#Software: Microsoft Internet Information Services 7.0
#Version: 1.0
#Date: 2019-11-12 22:22:22
#Fields: time c-ip cs-method cs-uri-stem sc-status sc-win32-status
22:22:22 175.45.176.200 [1]USER administrator 331 0
22:22:22 175.45.176.200 [2]USER administrator 331 0
22:22:22 175.45.176.200 [3]USER administrator 331 0
22:22:22 175.45.176.200 [4]USER administrator 331 0
22:22:22 175.45.176.200 [5]USER administrator 331 0
22:22:22 175.45.176.200 [2]PASS - 530 1326
22:22:22 175.45.176.200 [1]PASS - 530 1326
22:22:22 175.45.176.200 [3]PASS - 530 1326
22:22:22 175.45.176.200 [5]PASS - 530 1326
22:22:22 175.45.176.200 [2]USER administrator 331 0
22:22:22 175.45.176.200 [1]USER administrator 331 0
22:22:22 175.45.176.200 [3]USER administrator 331 0
22:22:22 175.45.176.200 [4]PASS - 530 1326
22:22:22 175.45.176.200 [5]USER administrator 331 0
22:22:22 175.45.176.200 [4]USER administrator 331 0
22:22:22 175.45.176.200 [2]PASS - 530 1326
22:22:22 175.45.176.200 [3]PASS - 530 1326
22:22:22 175.45.176.200 [2]USER administrator 331 0
22:22:22 175.45.176.200 [5]PASS - 530 1326
22:22:22 175.45.176.200 [4]PASS - 530 1326
22:22:22 175.45.176.200 [1]PASS - 530 1326
```

12

Type the following **command** and **press Enter** to view the **last 16 lines** of the **log files**.

```
root@kali:~# tail -n 16 ex191112.log
```



```
root@kali:~# tail -n 16 ex191112.log
22:25:42 175.45.176.200 [2]USER administrator 331 0
22:25:42 175.45.176.200 [5]PASS - 530 1326
22:25:42 175.45.176.200 [1]PASS - 530 1326
22:25:42 175.45.176.200 [5]USER administrator 331 0
22:25:42 175.45.176.200 [3]PASS - 530 1326
22:25:42 175.45.176.200 [1]USER administrator 331 0
22:25:42 175.45.176.200 [3]USER administrator 331 0
22:25:42 175.45.176.200 [5]PASS - 230 0
22:25:42 175.45.176.200 [2]PASS - 530 1326
22:25:42 175.45.176.200 [5]QUIT - 230 0
22:25:42 175.45.176.200 [2]USER administrator 331 0
22:25:42 175.45.176.200 [4]PASS - 530 1326
22:25:42 175.45.176.200 [3]PASS - 530 1326
22:25:42 175.45.176.200 [1]PASS - 530 1326
22:25:42 175.45.176.200 [2]PASS - 530 1326
```

The grep (global regular expressions print) command can be used to parse data for a certain term or search expression.

- 13 **Type** the following **command** and **press Enter** to view the options for **grep**. **grep** is a program that finds pattern matches in a file.

```
root@kali:~# grep --help
```

- 14 **Type** the following **command** and **press Enter** to view all lines with the **Word Microsoft**.

```
root@kali:~# cat ex191112.log | grep Microsoft
```

```
root@kali:~# cat ex191112.log | grep Microsoft
#Software: Microsoft Internet Information Services 7.0
```

- 15 **Type** the following **command** and **press Enter** to view all lines with the word **Date**.

```
root@kali:~# cat ex191112.log | grep Date
```

```
root@kali:~# cat ex191112.log | grep Date
#Date: 2019-11-12 22:22:22
```

The grep command is case sensitive. When we change the word date to all lowercase, there are no results.

- 16 **Type** the following **command** and **press Enter** to view all lines with the word **date**.

```
root@kali:~# cat ex191112.log | grep date
```

```
root@kali:~# cat ex191112.log | grep date
```



Note: There are no results.

- 17 **Type** the following **command** and **press Enter** to grep for the particular expression.

```
root@kali:~# cat ex191112.log | grep "PASS - 230"
```

```
root@kali:~# cat ex191112.log | grep "PASS - 230"
Binary file (standard input) matches
```

- 18 **Notice** the **error** you are receiving. This is due to the information at the beginning of the file. We will need to remove that information to remove the error. **Open** the **file** using **Leafpad**. Leafpad is an open-source text editor in Linux.

```
root@kali:~# leafpad ex191112.log
```



```
root@kali:~# leafpad ex191112.log
```

```
ex191112.log
File Edit Search Options Help
#Software: Microsoft Internet Information Services 7.0
#Version: 1.0
#Date: 2019-11-12 22:22:22
#Fields: time c-ip cs-method cs-uri-stem sc-status sc-win32-status
22:22:22 175.45.176.200 [1]USER administrator 331 0
22:22:22 175.45.176.200 [2]USER administrator 331 0
22:22:22 175.45.176.200 [3]USER administrator 331 0
22:22:22 175.45.176.200 [4]USER administrator 331 0
22:22:22 175.45.176.200 [5]USER administrator 331 0
22:22:22 175.45.176.200 [2]PASS - 530 1326
22:22:22 175.45.176.200 [1]PASS - 530 1326
22:22:22 175.45.176.200 [3]PASS - 530 1326
22:22:22 175.45.176.200 [5]PASS - 530 1326
22:22:22 175.45.176.200 [2]USER administrator 331 0
22:22:22 175.45.176.200 [1]USER administrator 331 0
22:22:22 175.45.176.200 [3]USER administrator 331 0
22:22:22 175.45.176.200 [4]PASS - 530 1326
22:22:22 175.45.176.200 [5]USER administrator 331 0
22:22:22 175.45.176.200 [4]USER administrator 331 0
22:22:22 175.45.176.200 [2]PASS - 530 1326
```

- 19 Remove the first four lines (as seen below) and then **save** the file. **Close** the log file.

```
ex191112.log
File Edit Search Options Help
22:22:22 175.45.176.200 [1]USER administrator 331 0
22:22:22 175.45.176.200 [2]USER administrator 331 0
22:22:22 175.45.176.200 [3]USER administrator 331 0
22:22:22 175.45.176.200 [4]USER administrator 331 0
22:22:22 175.45.176.200 [5]USER administrator 331 0
22:22:22 175.45.176.200 [2]PASS - 530 1326
22:22:22 175.45.176.200 [1]PASS - 530 1326
22:22:22 175.45.176.200 [3]PASS - 530 1326
22:22:22 175.45.176.200 [5]PASS - 530 1326
22:22:22 175.45.176.200 [2]USER administrator 331 0
22:22:22 175.45.176.200 [1]USER administrator 331 0
22:22:22 175.45.176.200 [3]USER administrator 331 0
22:22:22 175.45.176.200 [4]PASS - 530 1326
22:22:22 175.45.176.200 [5]USER administrator 331 0
22:22:22 175.45.176.200 [4]USER administrator 331 0
22:22:22 175.45.176.200 [2]PASS - 530 1326
22:22:22 175.45.176.200 [3]PASS - 530 1326
22:22:22 175.45.176.200 [2]USER administrator 331 0
22:22:22 175.45.176.200 [5]PASS - 530 1326
22:22:22 175.45.176.200 [4]PASS - 530 1326
```

- 20 Type the following **command** and **press Enter** to grep for the particular expression.

```
root@kali:~# cat ex191112.log | grep "PASS - 230"
```

```
root@kali:~# cat ex191112.log | grep "PASS - 230"  
22:25:42 175.45.176.200 [5]PASS - 230 0
```



Note: The command line expression `gawk` can be used to display a particular column of your output. The default delimiter is the space, but you can change that by typing the `gawk` command at the terminal to see all of the available options.

21

Type the following **command** and **press Enter** to view the options for the **gawk**

command. `gawk` is a program is a powerful pattern-matching and processing language.

```
root@kali:~# gawk --help
```

22

Type the following **command** and **press Enter** to use the `gawk` command and display

the **first column** in the **log file**.

```
root@kali:~# cat ex191112.log | gawk '{print $1}'
```


25

Type the following **command** and **press Enter** to use the gawk command and display the total number of each of the **unique values** in the **first column** in the **log file**.

```
root@kali:~# cat ex191112.log | gawk '{print $1}' | sort | uniq -c
```

```
root@kali:~# cat ex191112.log | gawk '{ print $1}' | sort | uniq -c
1835 22:22:22
1812 22:22:23
1902 22:22:24
1966 22:22:25
1957 22:22:26
1877 22:22:27
1960 22:22:28
1952 22:22:29
1993 22:22:30
1981 22:22:31
1897 22:22:32
1926 22:22:33
1942 22:22:34
1977 22:22:35
2000 22:22:36
1991 22:22:37
1822 22:22:38
1963 22:22:39
1969 22:22:40
1974 22:22:41
1938 22:22:42
1944 22:22:43
```

26

Type the following **command** and **press Enter** to use the gawk command and display the total number of each of the **unique values** in the **three columns** in the **log file**. In Linux, you can send the output of a command to another command using a **|** character.

```
root@kali:~# cat ex191112.log | gawk '{print $3,$4,$5}' | sort | uniq -c
```

```
root@kali:~# cat ex191112.log | gawk '{ print $3, $4, $5}' | sort | uniq -c
36176 [1]PASS - 530
36176 [1]USER administrator 331
36101 [2]PASS - 530
36101 [2]USER administrator 331
36185 [3]PASS - 530
36185 [3]USER administrator 331
36165 [4]PASS - 530
36165 [4]USER administrator 331
    1 [5]PASS - 230
36228 [5]PASS - 530
    1 [5]QUIT - 230
36229 [5]USER administrator 331
```

FTP log codes include:

- 331 – attempted log in
- 530 – failure to log in
- 230 – successful login/logout

Here is a reference: https://en.wikipedia.org/wiki/List_of_FTP_server_return_codes

27 **Type** the following **command** and **press Enter** to use the gawk command and display the total number of each of the **unique values** in the **fifth column** in the **log file**.

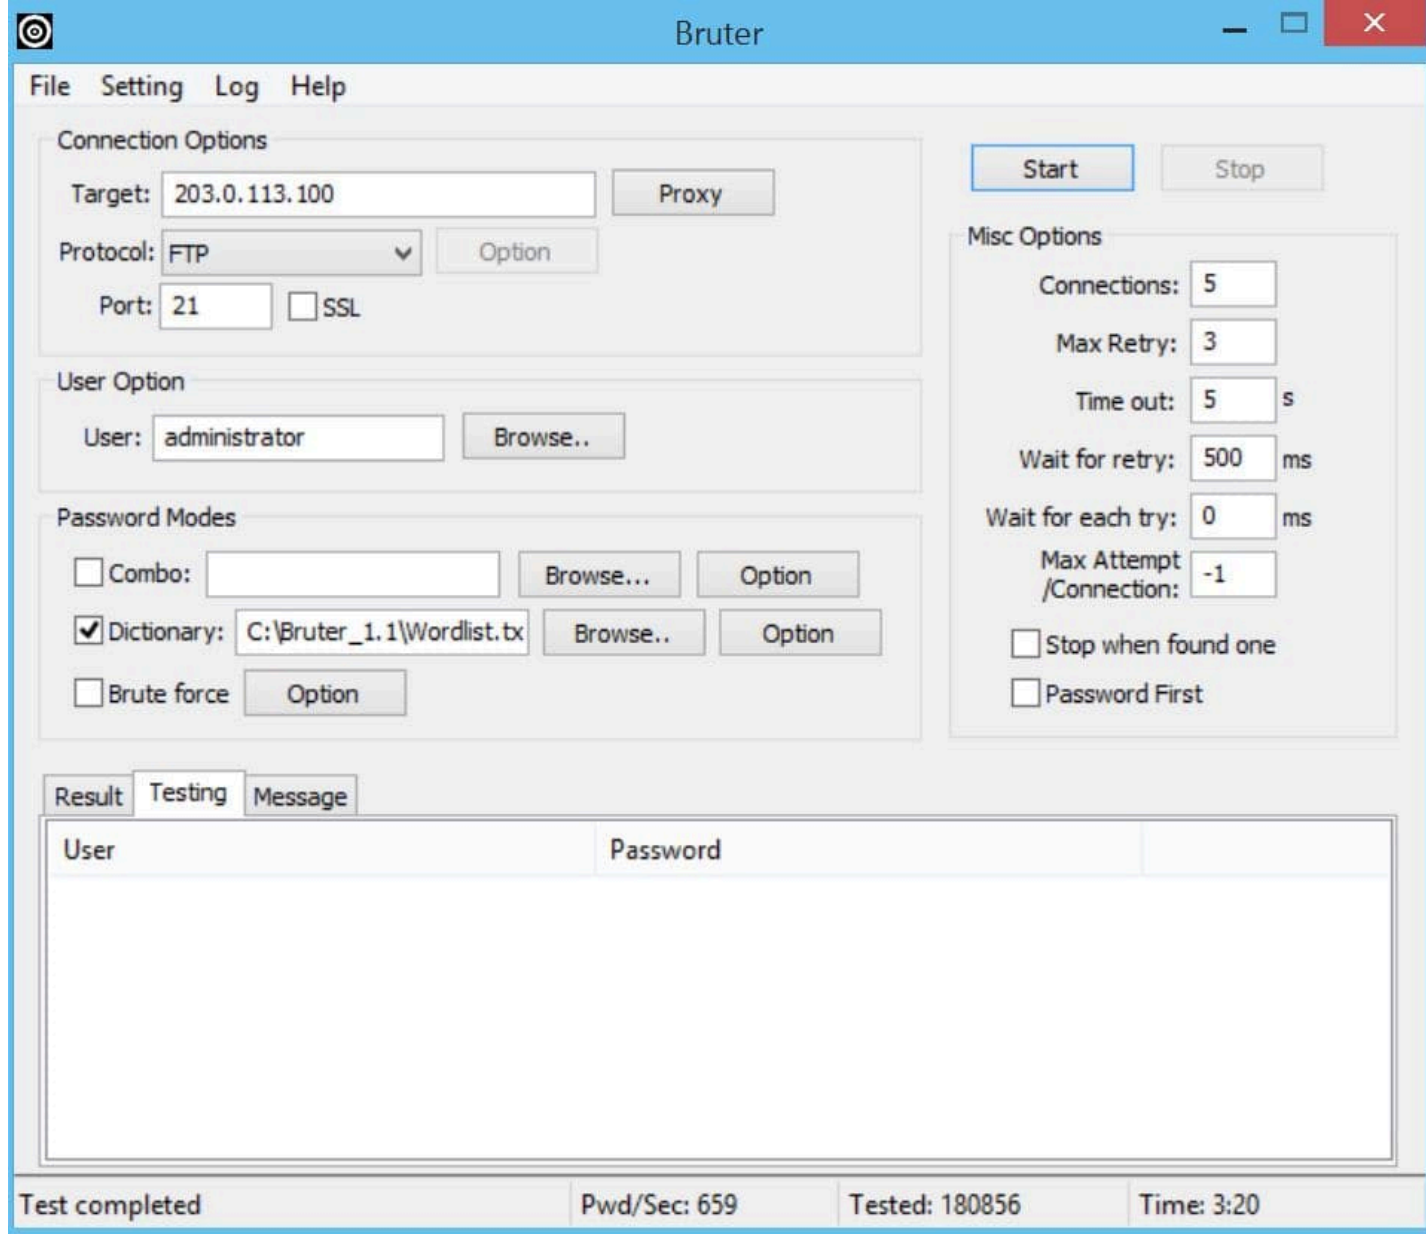
```
root@kali:~# cat ex191112.log | gawk '{print $5}' | sort | uniq -c
```

```
root@kali:~# cat ex191112.log | gawk '{ print $5}' | sort | uniq -c
    2 230
180856 331
180855 530
```

According to the attacker, there were 180,856 attempted logins by the administrator account.

The 331 code shows this to be the case.

There were 180,855 failed logins shown by the error code of 530.



Summary

Log files can be extremely large. One way to view their output is by using the command line in Linux. The cat command in Linux will allow you to view a file, and the grep command can be used to parse through the information. The head command will show you the first lines in a file, and the tail command will show you the last few lines. Using a combination of the gawk, sort, and uniq commands, you can find out specific information about events in log files.

DISCUSSION QUESTIONS:

- 1 What is the command to view a file in Linux?
- 2 What is the command to count the number of lines in a file in Linux?
- 3 What is the default delimiter in gawk?
- 4 What is the command to rename a file in Linux?

5

What is the command to print a column in Linux?

Splunk

Splunk is a log analysis tool that will allow you to ingest several different types of log formats, including web logs, system logs, event viewer files, and PCAP data. Splunk is a commercial product, but it has a free edition. With the free edition of Splunk, there is a parse limit of 500 MB per day within a 24-hour time frame.

1

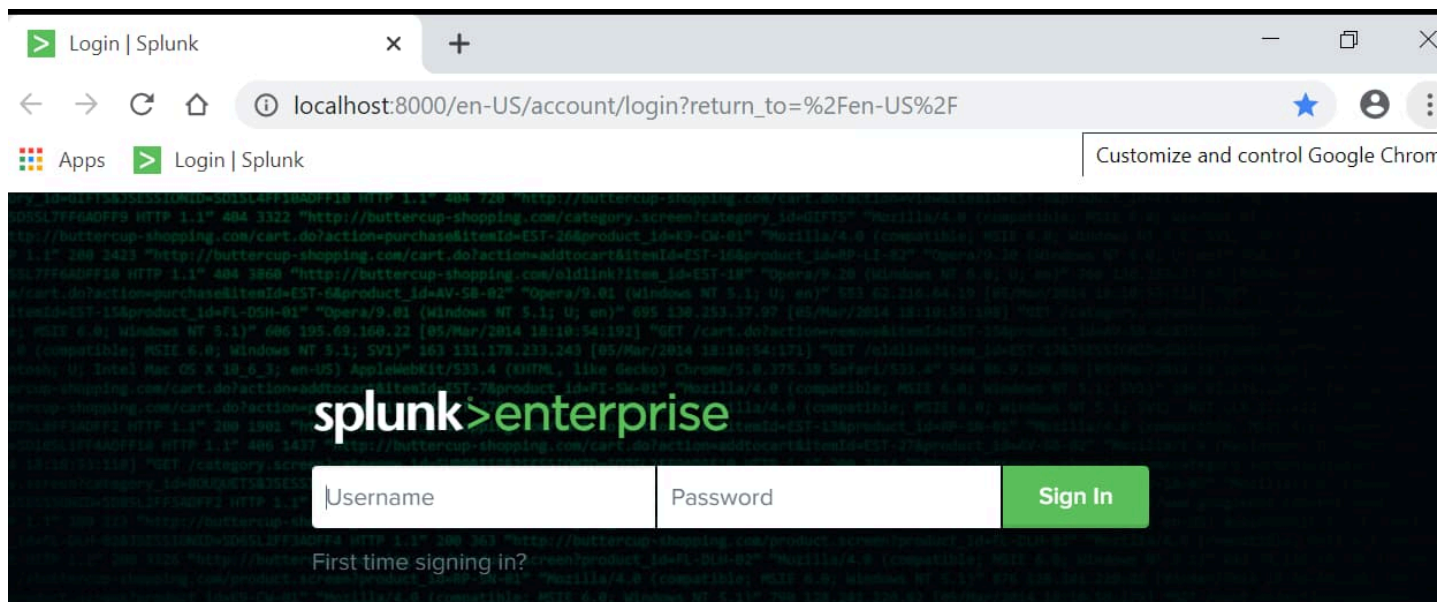
Log on to the **Windows 10 VM** as **Student** with the password of **P@ssw0rd**. **Double-click** the **Chrome icon** on the desktop.



Note: The password of **P@ssw0rd** will not be displayed when you type it for security purposes.


2

You will automatically be logged in to **Splunk**.



3

Click **Explore Splunk** in the top right of the window.



4

Click Add Data. **Click** Skip Tour if asked.



Add Data

Add or forward data to Splunk Enterprise. Afterwards, you may [extract fields](#).

5

Ignore any **errors** about not being connected to the **Internet**. **Click** Upload files from my computer.



Upload files from my computer

Local log files
Local structured files (e.g. CSV)
[Tutorial for adding data](#)

6

Click **Select File**.

splunk>enterprise Apps Admin Messages Settings Activity Help Find

Add Data ● ○ ○ ○ ○ < Back Next >

Select Source Set Source Type Input Settings Review Done

Select Source

Choose a file to upload to the Splunk platform, either by browsing your computer or by dropping a file into the target box below. [Learn More](#)

Selected File: **No file selected**

Select File

Drop your data file here

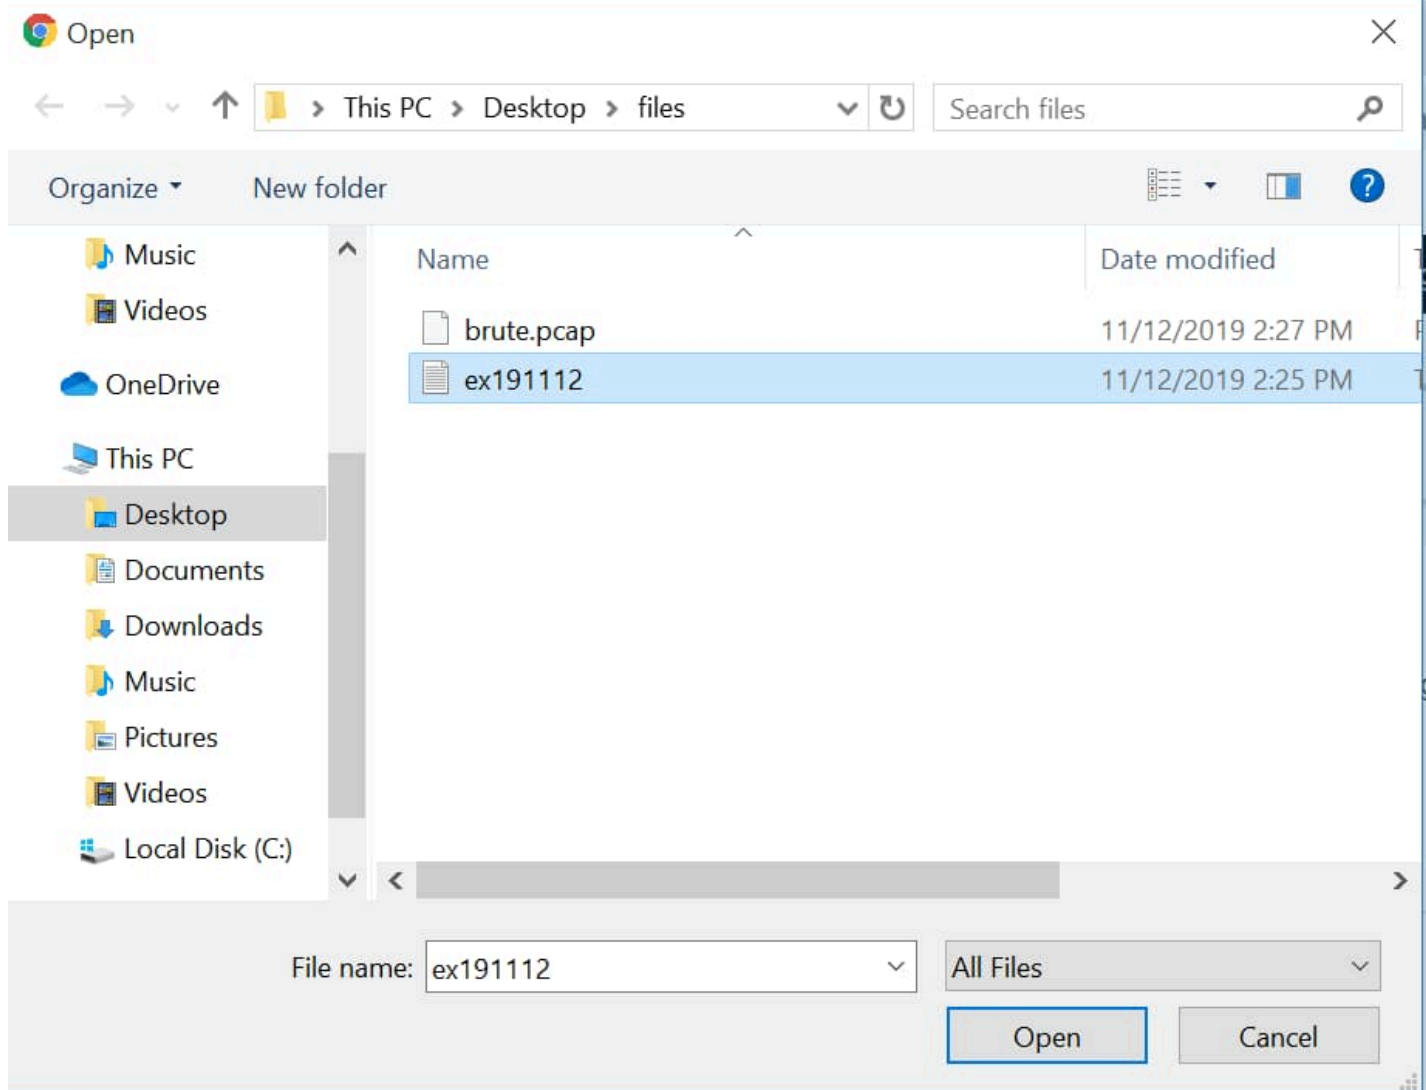
The maximum file upload size is 500 Mb



Note: There is a limit of 500 MB per day of data with the free version of Splunk.

7

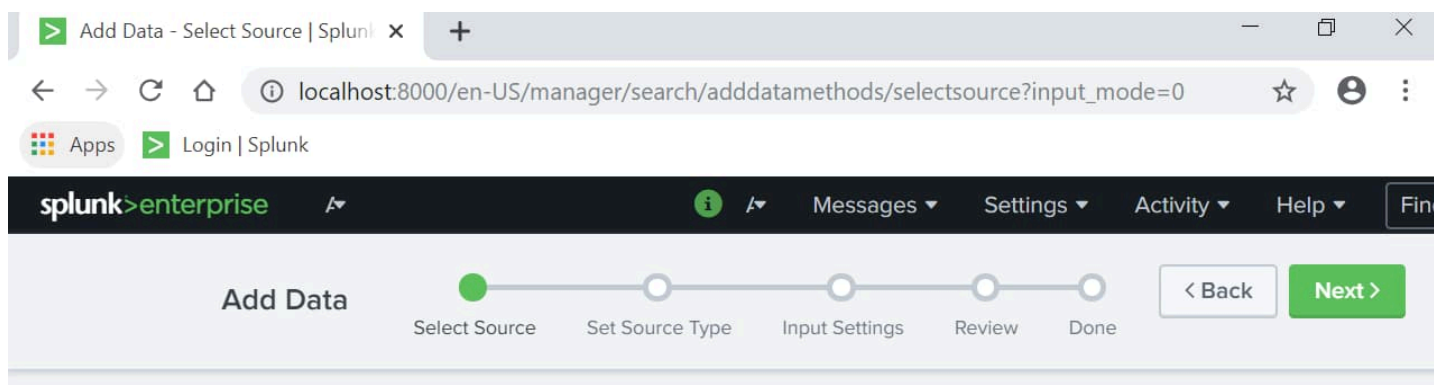
Browse to **This PC > Desktop > files**. **Double-click** on the **ex19112** file.



Note: This is the same file we parsed through with the Linux command line tools.

8

Click **Next**.



9

Verify that the **Source type** is **iis** and **select Next**.

Browser tabs: Add Data - Set Sourcetype | Splunk x +

Address bar: localhost:8000/en-US/manager/search/adddatamethods/datapreview

Navigation: Apps Login | Splunk

Splunk > enterprise Messages Settings Activity Help Find

Add Data Progress: Select Source (Active) Set Source Type Input Settings Review Done

< Back Next >

Set Source Type

This page lets you see how the Splunk platform sees your data before indexing. If the events look correct and have the right timestamps "Next" to proceed. If not, use the options below to define proper event breaks and timestamps. If you cannot find an appropriate source type for your data, create a new one by clicking "Save As".

Source: **ex191112.log** [View Event S](#)

Source type: iis Save As Table Format 20 Per Page



NOTE: The source type is the type of log file that Splunk is going to analyze. iis is a web server on Windows web servers.

10

Click **Review** to **Input Settings**.

Browser tabs: Add Data - Input Settings | Splunk x +

Address bar: localhost:8000/en-US/manager/search/adddatamethods/inputsettings

Navigation: Apps Login | Splunk

Splunk > enterprise Messages Settings Activity Help Find

Add Data Progress: Select Source Set Source Type (Active) Input Settings Review Done

< Back Review >

Input Settings

Optionally set additional input parameters for this data input as follows:

11

Click **Submit**.

Browser window: Add Data - Review | Splunk 8.0.0

URL: localhost:8000/en-US/manager/search/adddatamethods/review

Navigation: Apps, Login | Splunk

splunk>enterprise

Messages, Settings, Activity, Help

Add Data

Progress: Select Source, Set Source Type, Input Settings, Review, Done

< Back Submit >

Review

Input Type Uploaded File
File Name ex191112.log
Source Type iis
Host DESKTOP-A2KB24I
Index Default

12

Click **Start Searching**.

Browser window: Add Data - Success | Splunk 8.0.0

URL: localhost:8000/en-US/manager/search/adddatamethods/success

Navigation: Apps, Login | Splunk

splunk>enterprise

Messages, Settings, Activity, Help

Add Data

Progress: Select Source, Set Source Type, Input Settings, Review, Done

< Back Next >

✓ File has been uploaded successfully.
Configure your inputs by going to Settings > [Data Inputs](#)

Start Searching Search your data now or see [examples and tutorials](#).

13

In the **New Search bar**, **type** `"530"` and then **click** the **search button**. The "530" event code in IIS is a User cannot login error.

Search | Splunk 8.0.0

localhost:8000/en-US/app/search/search?q=search%20"530"&earliest=0&latest=&d...

Apps Login | Splunk

New Search

Save As Close

"530" All time

✓ 180,855 events (before 11/17/19 7:01:36.000 PM) No Event Sampling Job

Events (180,855) Patterns Statistics Visualization

Format Timeline Zoom Out Zoom to Selection Deselect 1 minute per column

List Format 20 Per Page

< Prev 1 2 3 4 5 6 7 8 ... Next >

< Hide Fields All Fields

SELECTED FIELDS

- host 1
- source 1
- sourcetype 1

INTERESTING FIELDS

- c_ip 1
- cs_method 5

i	Time	Event
>	11/12/19 2:25:42.000 PM	22:25:42 175.45.176.200 [2]PASS - 530 1326 host = DESKTOP-A2KB24I source = ex191112.log sourcetype = iis
>	11/12/19 2:25:42.000 PM	22:25:42 175.45.176.200 [1]PASS - 530 1326 host = DESKTOP-A2KB24I source = ex191112.log sourcetype = iis
>	11/12/19 2:25:42.000 PM	22:25:42 175.45.176.200 [3]PASS - 530 1326 host = DESKTOP-A2KB24I source = ex191112.log sourcetype = iis

The number of events with the code of 530 is 180,855. This was the exact same number of these event codes when we used all of the Linux commands below to parse the information.

```
root@kali:~# cat ex191112.log | gawk '{print $5}' | sort | uniq -c
```

```
root@kali:~# cat ex191112.log | gawk '{ print $5}' | sort | uniq -c
  2 230
180856 331
180855 530
```

Notice that both methods provide the same result, but the Linux results were provided faster. The disadvantage to the Linux method is remembering the use and syntax of the commands.

14

In the **New Search bar**, type **"230"** and then **click** the **search button**.

"230" event code is a successfully log in in IIS.

The screenshot shows the Splunk Search interface. The search bar contains the query "230". The results bar indicates "2 events (before 11/17/19 6:58:03.000 PM)". The results table shows two events, both with the code 230.

Time	Event
11/12/19 2:25:42.000 PM	22:25:42 175.45.176.200 [5]QUIT - 230 0
11/12/19 2:25:42.000 PM	22:25:42 175.45.176.200 [5]PASS - 230 0

The number of events with the code of 230 is 2. This was the exact same number of these event codes when we used all of the Linux commands below to parse the information.

```
root@kali:~# cat ex191112.log | gawk '{print $5}' | sort | uniq -c
```

The terminal window shows the output of the command `cat ex191112.log | gawk '{print $5}' | sort | uniq -c`. The output is:

```
2 230
180856 331
180855 530
```

A red box highlights the first line of output, "2 230", and a red arrow points to it from the left.

Again, both methods provide the same result, but the Linux results were provided faster. As we started before, the major disadvantage to the Linux method is remembering the use and syntax of the commands. Frequent use and practice with the Linux commands and the available switches will make you more familiar with them. Using Splunk is always another option.



Note: **Press** the **STOP** button to complete the lab.

Summary

Splunk is a log analysis tool that will allow you to ingest several different types of log formats, including web logs, system logs, event viewer files, and PCAP data. Splunk is a commercial product, but it has a free edition. With the free edition of Splunk, there is a parse limit of 500 MB per day within a 24-hour time frame. Splunk and Linux command line log analysis provide the same result, but the Linux results are provided faster. The disadvantage to the Linux method is remembering the use and syntax of the commands. Frequent use and practice with the Linux commands and the available switches will make you more familiar with them. Using Splunk is always another option.

DISCUSSION QUESTIONS:

- 1 What is Splunk?
- 2 What are some of the advantages of using Splunk?
- 3 What are some advantages of using the Linux command line over Splunk?
- 4 Is Splunk free?
- 5 What limitation will you face if you do not pay for Splunk?